

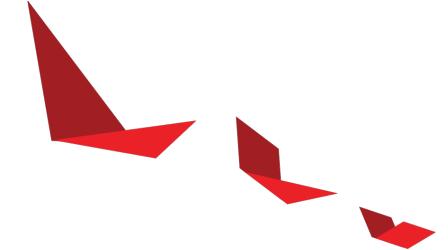
ROS 2

# Hardware Acceleration Working Group

**Meeting #2:** ROS 2 architecture for hardware acceleration, use case and ReconROS



# Agenda



Welcome and newcomers' intro



Recap of ROS 2  
Hardware Acceleration  
WG progress



Adding new boards to the  
HAWG architecture progress



Use case



ReconROS (invited  
talk)



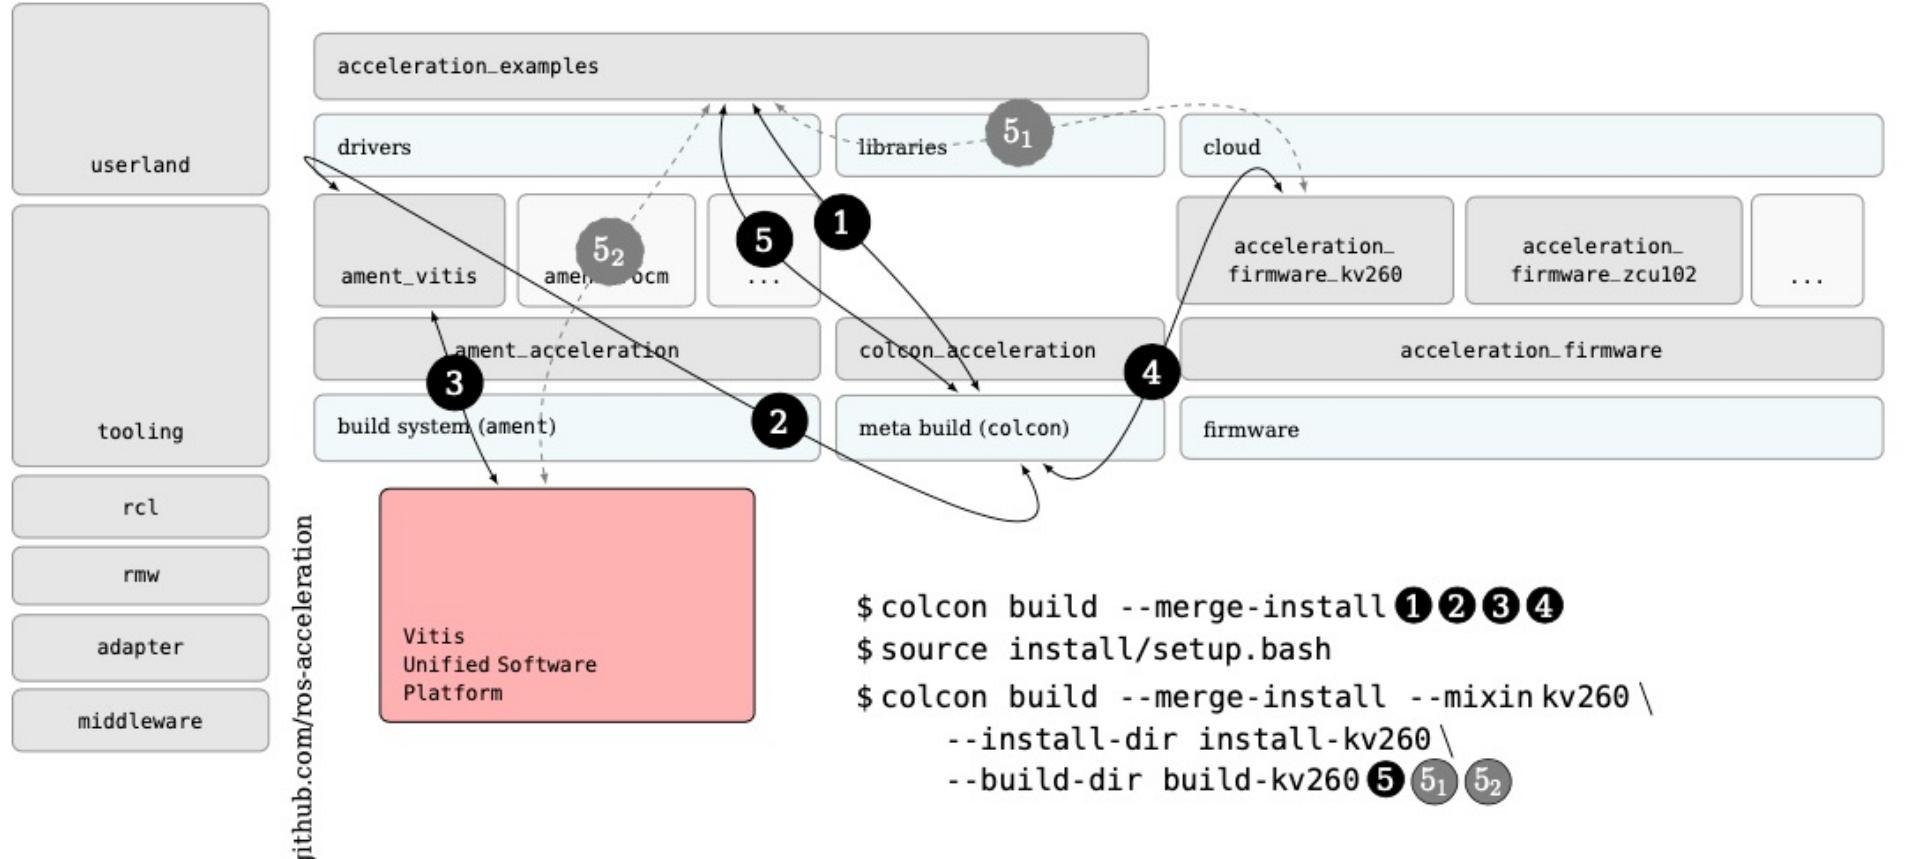
Q&A

# Recap of ROS 2 Hardware Acceleration WG resources

**Objective:** Drive creation, maintenance and testing of acceleration kernels on top of open standards (C++ and OpenCL) for optimized ROS 2 and Gazebo interactions over different compute substrates (including FPGAs, GPUs and ASICs).

- **Reference hardware platform:** [Kria K26 Adaptive SOM](#)
- **Meeting invite group:** [ROS 2 Hardware Acceleration WG Google Group](#)
- **Instant messaging:** [Matrix community](#) (Matrix is an open network for secure, decentralized communication).
- **Backlog management:** [Phase 1: tools, examples, benchmarking and first demonstrators](#)
- **Github organization:** [ros-acceleration](#)
- Meeting minutes: [HAWG minutes](#)
- **Discourse tag:** [wg-acceleration](#)

# Recap of ROS 2 Hardware Acceleration WG architecture



For more on architecture:

<https://github.com/ros-infrastructure/rep/pull/324>

Open source component

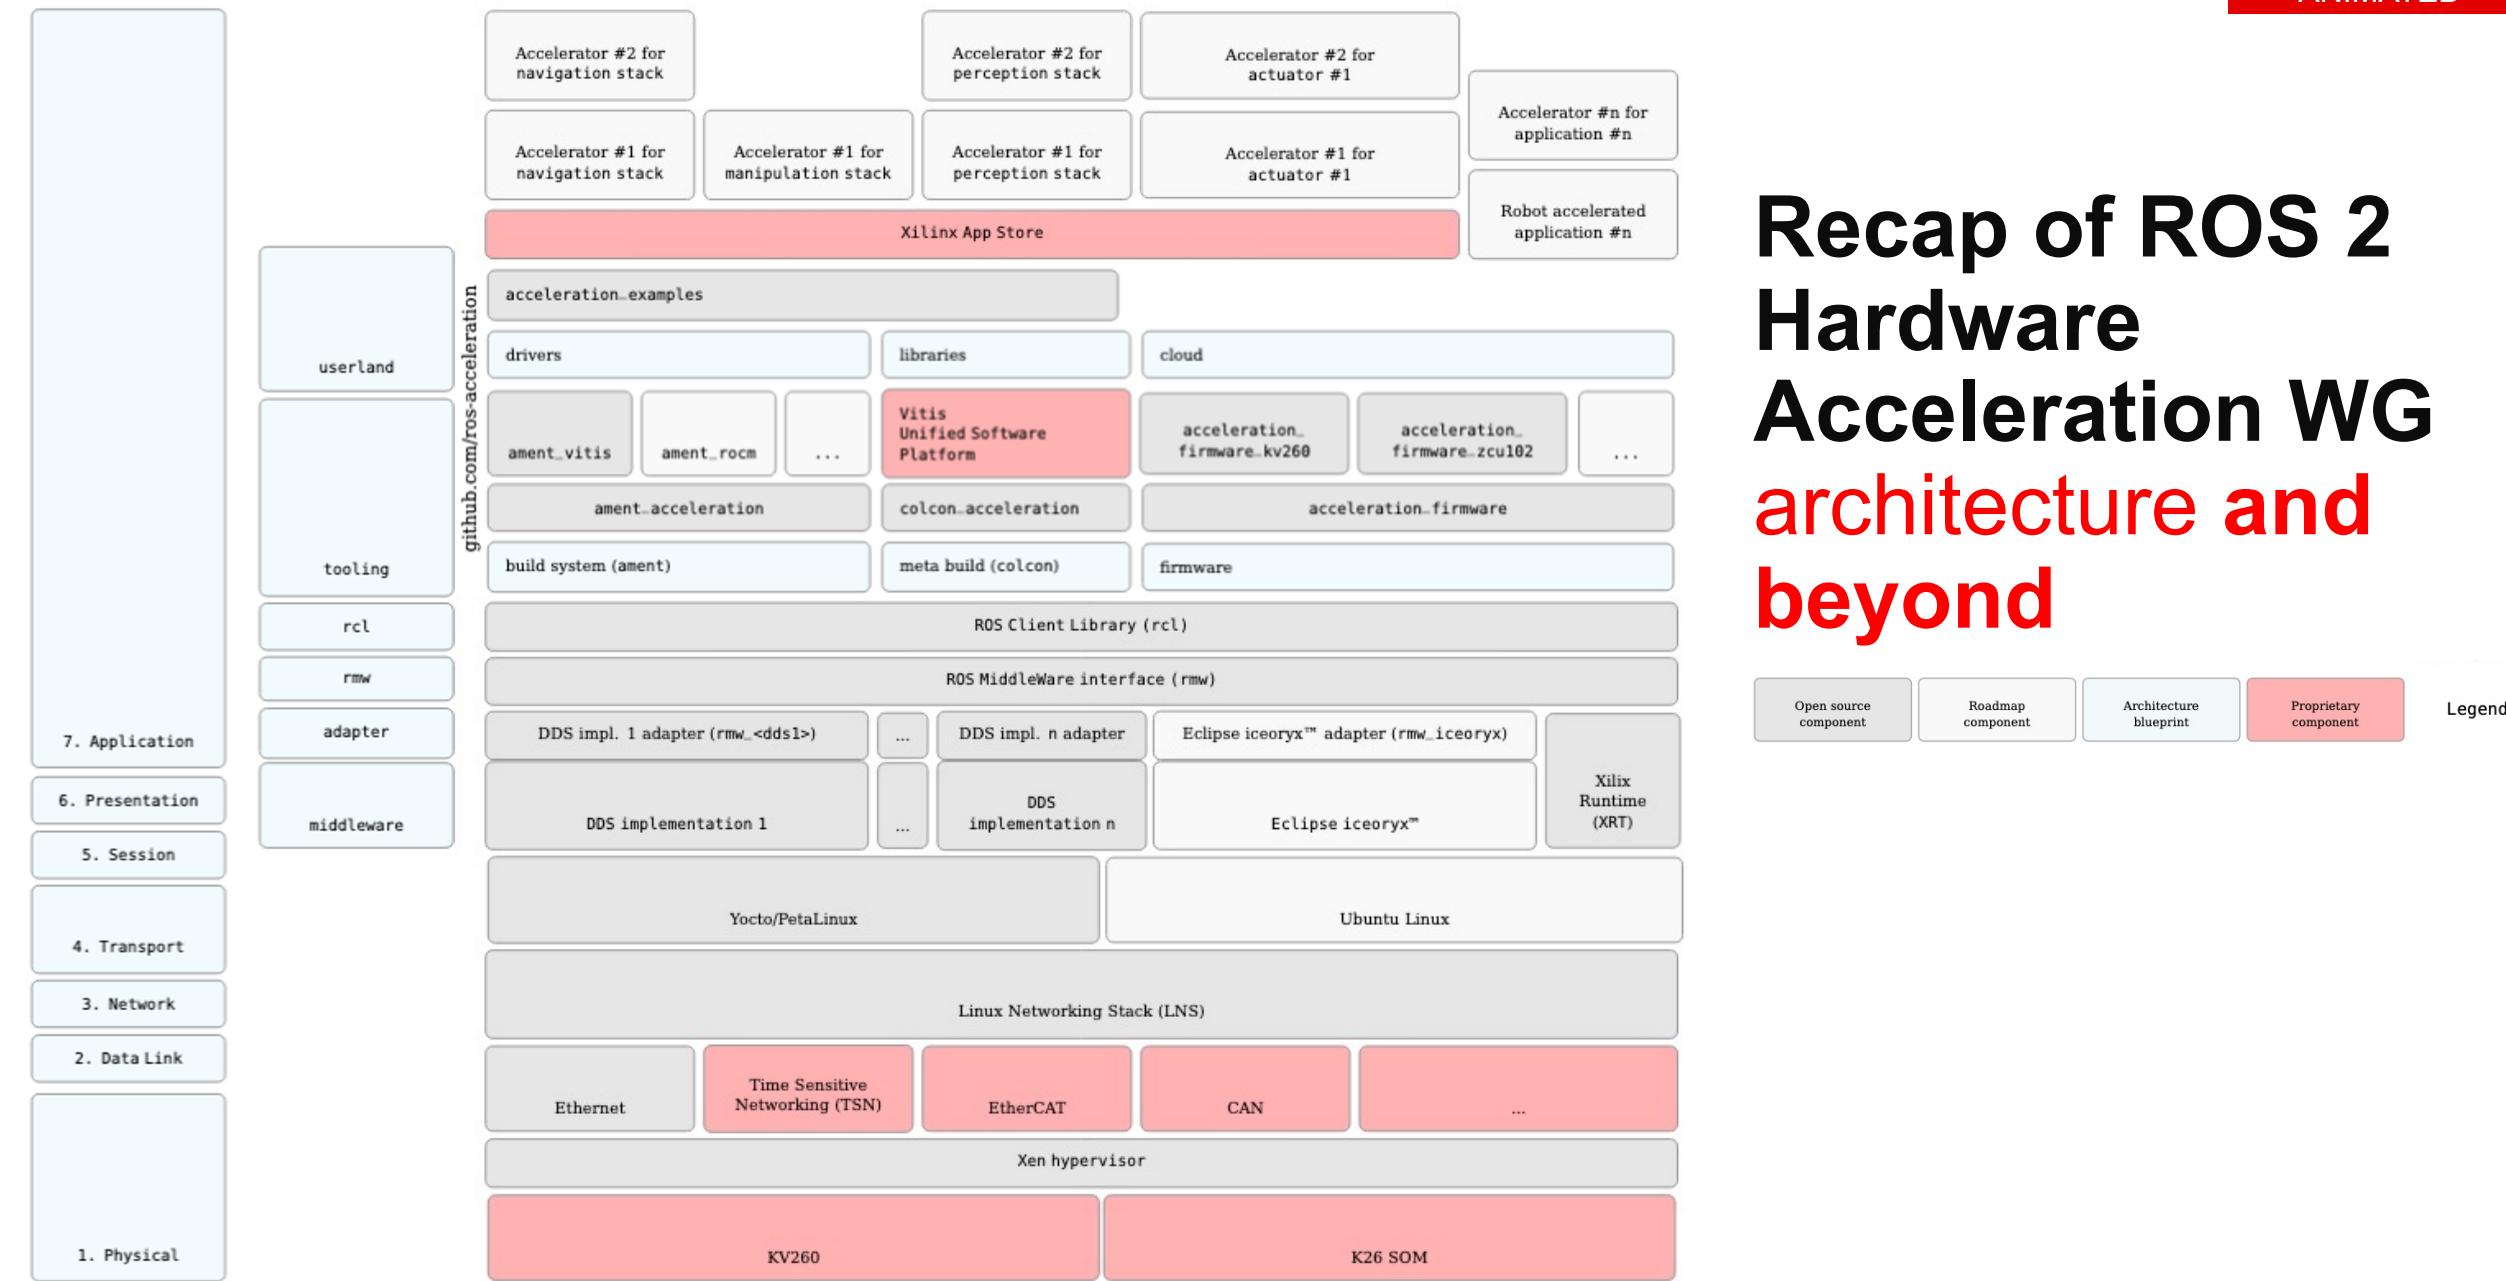
Roadmap component

Architecture blueprint

Proprietary component

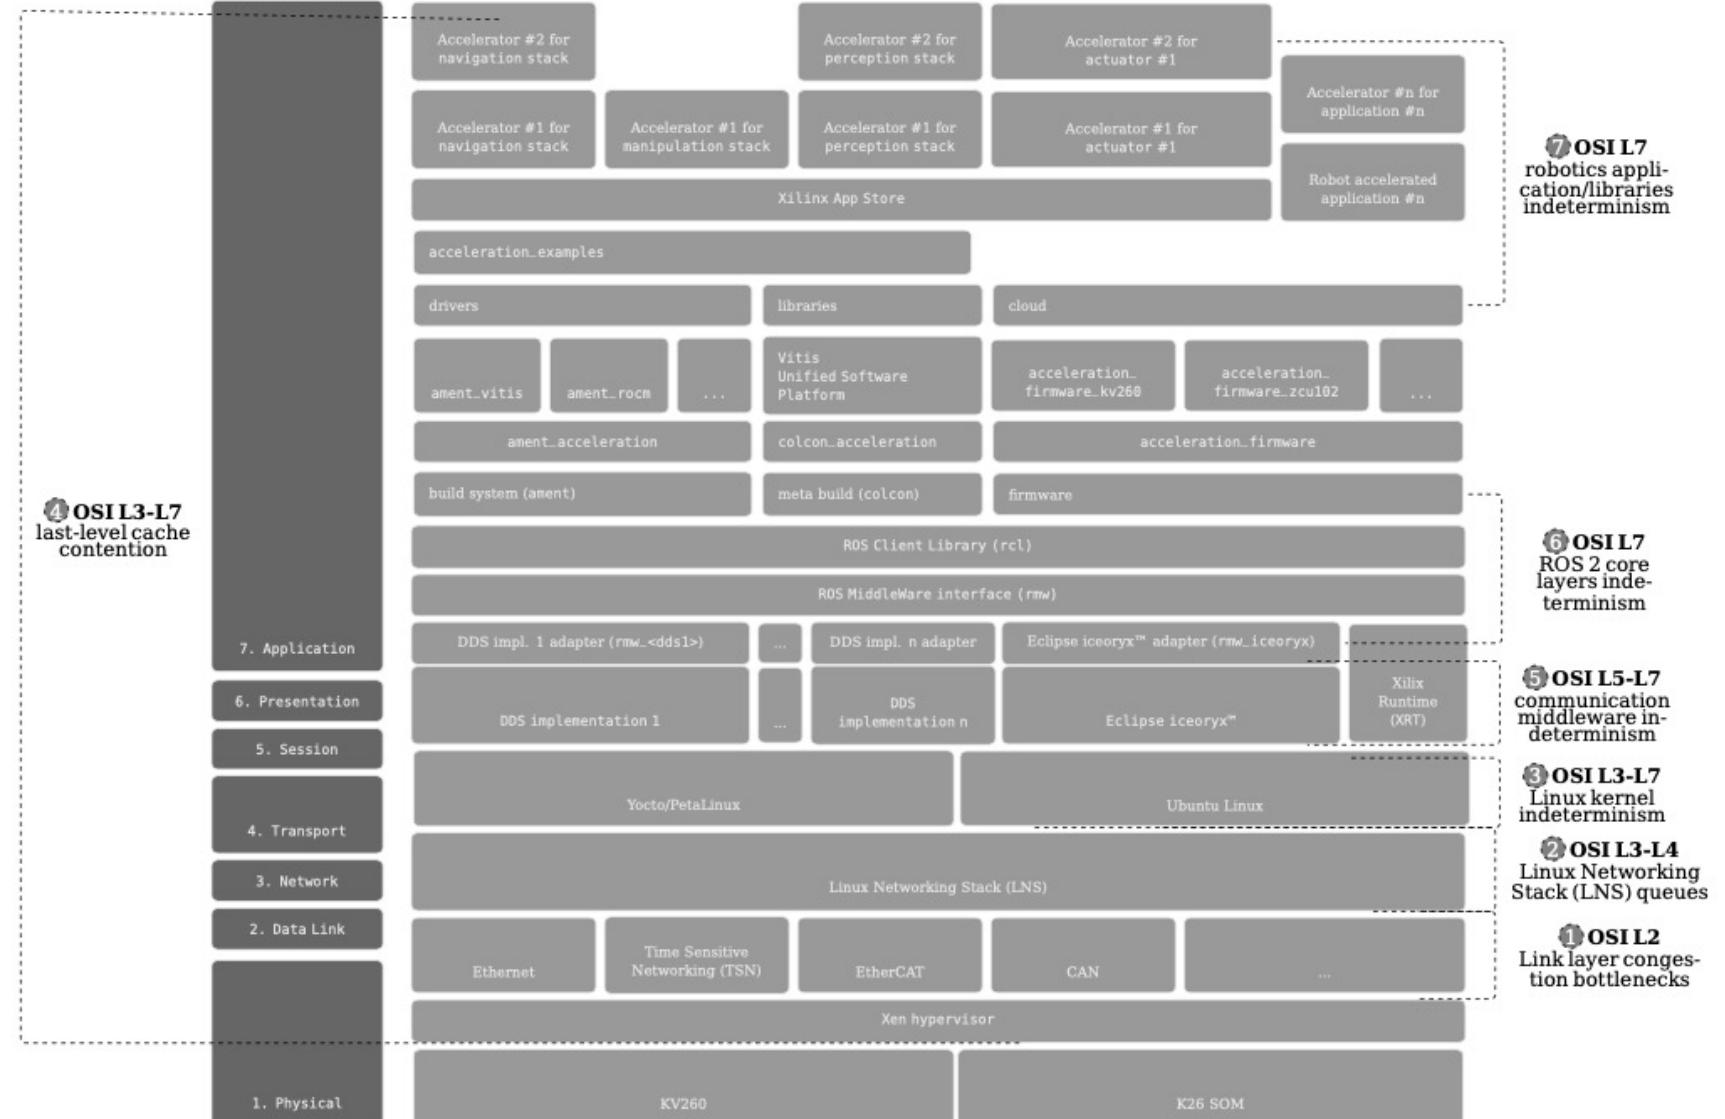
Legend

# Recap of ROS 2 Hardware Acceleration WG architecture and beyond



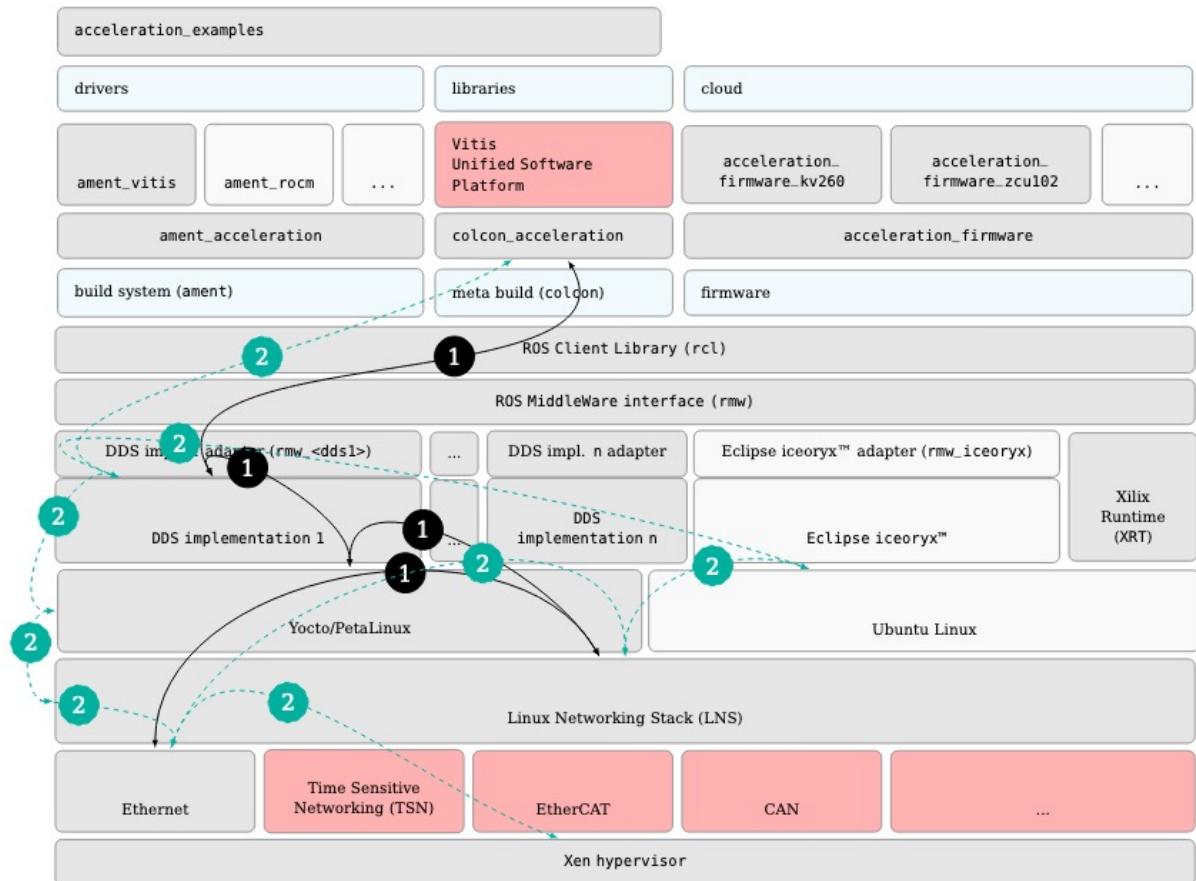
# Delivering real-time ROS 2

An end-to-end approach from the silicon layer to mitigate all sources of indeterminism across the ROS robotics stack.



# Xilinx facilitates creating ROS 2 mixed critical systems

- 1 Command builds a raw image with the default ROS 2 configuration, a fully preemptible kernel (`preempt_rt`), the default PetaLinux-based rootfs, the LNS and an Ethernet link layer.
- 2 Command builds a raw image leveraging the Xen hypervisor with 3 VMs. The first, `dom0`, uses a vanilla kernel. The second, `domU1`, uses a vanilla kernel and a custom Ubuntu-based file system. The third is a `dom0less` VM and uses a fully preemptible kernel (`preempt_rt`). Unless otherwise specified, all VMs use the default ROS 2 configuration, PetaLinux-based rootfs, the LNS and an Ethernet link layer.



```
$ colcon acceleration kernel preempt_rt ①
```

Command builds a raw image with the default ROS 2 configuration, a fully preemptible kernel (`preempt_rt`), the default PetaLinux-based rootfs, the LNS and an Ethernet link layer.

```
$ colcon acceleration hypervisor \
  --dom0 vanilla \
  --domU vanilla -rootfs ubuntu.cpio.gz \
  --dom0less preempt_rt ②
```

Command builds a raw image leveraging the Xen hypervisor with 3 VMs. The first, `dom0`, uses a vanilla kernel. The second, `domU1`, uses a vanilla kernel and a custom Ubuntu-based file system. The third is a `dom0less` VM and uses a fully preemptible kernel (`preempt_rt`). Unless otherwise specified, all VMs use the default ROS 2 configuration, PetaLinux-based rootfs, the LNS and an Ethernet link layer.

Open source component

Roadmap component

Architecture blueprint

Proprietary component

Legend

# Kria Robotics Stack

## Bringing ROS 2 down to silicon

Empowering Hardware Acceleration and next-gen “robot chips”



# Kria Robotics Stack (KRS)

**Bringing ROS 2  
down to silicon**

Empowering  
Hardware Acceleration and  
next-gen “robot chips”



```
# 0. install Vitis 2020.2.2 and ROS Foxy

# 1. install some dependencies you might be missing
sudo apt-get -y install curl build-essential libssl-dev git wget
    \ ocl-icd-* opencl-headers python3-vcstool
    \ python3-colcon-common-extensions python3-colcon-mixin
    \ kpartx u-boot-tools

# 2. create a new ROS 2 workspace
mkdir -p ~/krs_ws/src; cd ~/krs_ws

# 3. KRS alpha release
cat << 'EOF' > krs_alpha.repos
repositories:
    acceleration/acceleration_firmware:
        type: git
        url: https://github.com/ros-acceleration/acceleration_firmware
        version: 0.4.0
    acceleration/acceleration_firmware_kv260:
        type: zip
        url: https://github.com/ros-acceleration/acceleration_firmware_kv260
            /releases/download/v0.6.0/acceleration_firmware_kv260.zip
    acceleration/colcon-acceleration:
        type: git
        url: https://github.com/ros-acceleration/colcon-acceleration
        version: 0.1.0
    acceleration/ros2acceleration:
        type: git
        url: https://github.com/ros-acceleration/ros2acceleration
        version: 0.2.0
    acceleration/ament_vitis:
        type: git
        url: https://github.com/ros-acceleration/ament_vitis
        version: 0.5.0
EOF

# 4. import repos of KRS alpha release
vcs import src --recursive < krs_alpha.repos # about 5 mins

# 5. build the workspace and deploy firmware for hardware acceleration
source /tools/Xilinx/Vitis/2020.2/settings64.sh ; source /opt/ros/foxy/setup.bash \
    ; export PATH="/usr/bin":$PATH
colcon build --merge-install # about 2 mins

# 6. source the overlay to enable all hardware acceleration features
source install/setup.bash
```

# Learn more about KRS

## Adaptive Computing in Robotics



WP537 (v1.0)

### Adaptive Computing in Robotics Leveraging ROS 2 to Enable Software-Defined Hardware for FPGAs

By: Víctor Mayoral-Vilches and Giulio Corradi

Xilinx's adaptive SOMs are the perfect compute platform for robotics, allowing the creation of software-defined hardware for robots, delivering solutions with increased performance per watt while also being cost-effective, energy efficient, secure, safe, and adaptable.

[xilinx.com/krs](https://www.xilinx.com/adaptive-computing-in-robotics.html)

## The Kria Robotics Stack (KRS)



WP540 (v1.0 DRAFT)

### Kria Robotics Stack

A ROS 2-centric Approach for Hardware  
Acceleration in Robotics

By: Victor Mayoral-Vilches

Hardware acceleration is revolutionizing robotics—empowering robots with the ability to react faster, consume less power, and present more secure architectures. Enter the future of robot chips built with ROS 2 as its SDK and specialized robotic circuitry built with SoCs.

[xilinx.com/krs](https://www.xilinx.com/kria-robotics-stack.html)

Access KRS alpha release [here](#) (documentation).  
[EA of KRS beta](#)

# Adding new boards to the HAWG – process

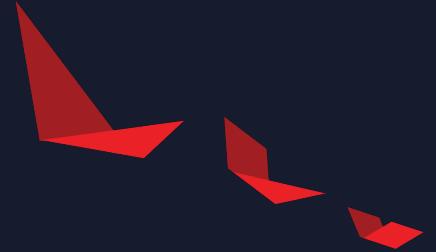
1. Check out the [Initial draft of REP-2008 - ROS 2 Hardware Acceleration Architecture and Conventions](#)
2. Create your own firmware repository for the corresponding board (see [acceleration firmware kv260](#) for an example)
3. Assess the capabilities of the hardware according to REP-2008 and create a table in the README.md that argues about it (see [example here](#))
4. Submit a PR to [ros-acceleration/community](#) to add your board to the community, with the corresponding support level according to REP-2008



## Ultra96-v2

Led by [@pimartos](#), check it out at

<https://github.com/ros-acceleration/community/issues/1>



# ROS 2 hardware acceleration use case

## robotics architect creating an accelerator for a faster ROS 2 publisher

# Accelerating ROS 2

## ROS 2 Hardware acceleration “hello world” – doublevadd\_publisher<sup>1</sup>

<sup>1</sup> [https://github.com/ros-acceleration/acceleration\\_examples/tree/main/doublevadd\\_publisher](https://github.com/ros-acceleration/acceleration_examples/tree/main/doublevadd_publisher)

```
int main(int argc, char * argv[]) {
    rclcpp::init(argc, argv);
    auto node = rclcpp::Node::make_shared("doublevadd_publisher");
    auto publisher = node->create_publisher<
        std_msgs::msg::String>("vector", 10);
    auto publish_count = 0;
    std_msgs::msg::String message;
    rclcpp::WallRate loop_rate(100ms); Goal: 10 Hz
    ...
    while (rclcpp::ok()) {
        ...
        vadd(in1, in2, out, DATA_SIZE);
        check_vadd(in1, in2, out);
        ...
        try {
            publisher->publish(message);
            rclcpp::spin_some(node);
        } catch (const rclcpp::exceptions::RCLError & e) {
            RCLCPP_ERROR(
                node->get_logger(),
                "unexpectedly failed with %s",
                e.what());
        }
        loop_rate.sleep();
    }
    rclcpp::shutdown();
    return 0;
}
```

Leo “the roboticist”



```
void vadd(
    const unsigned int *in1, // Read-Only Vector 1
    const unsigned int *in2, // Read-Only Vector 2
    unsigned int *out, // Output Result
    int size // Size in integer
)
{
    for (int j = 0; j < size; ++j) { // stupidly iterate
        // to generate "double" load
        for (int i = 0; i < size; ++i) {
            out[i] = in1[i] + in2[i];
        }
    }
}
```

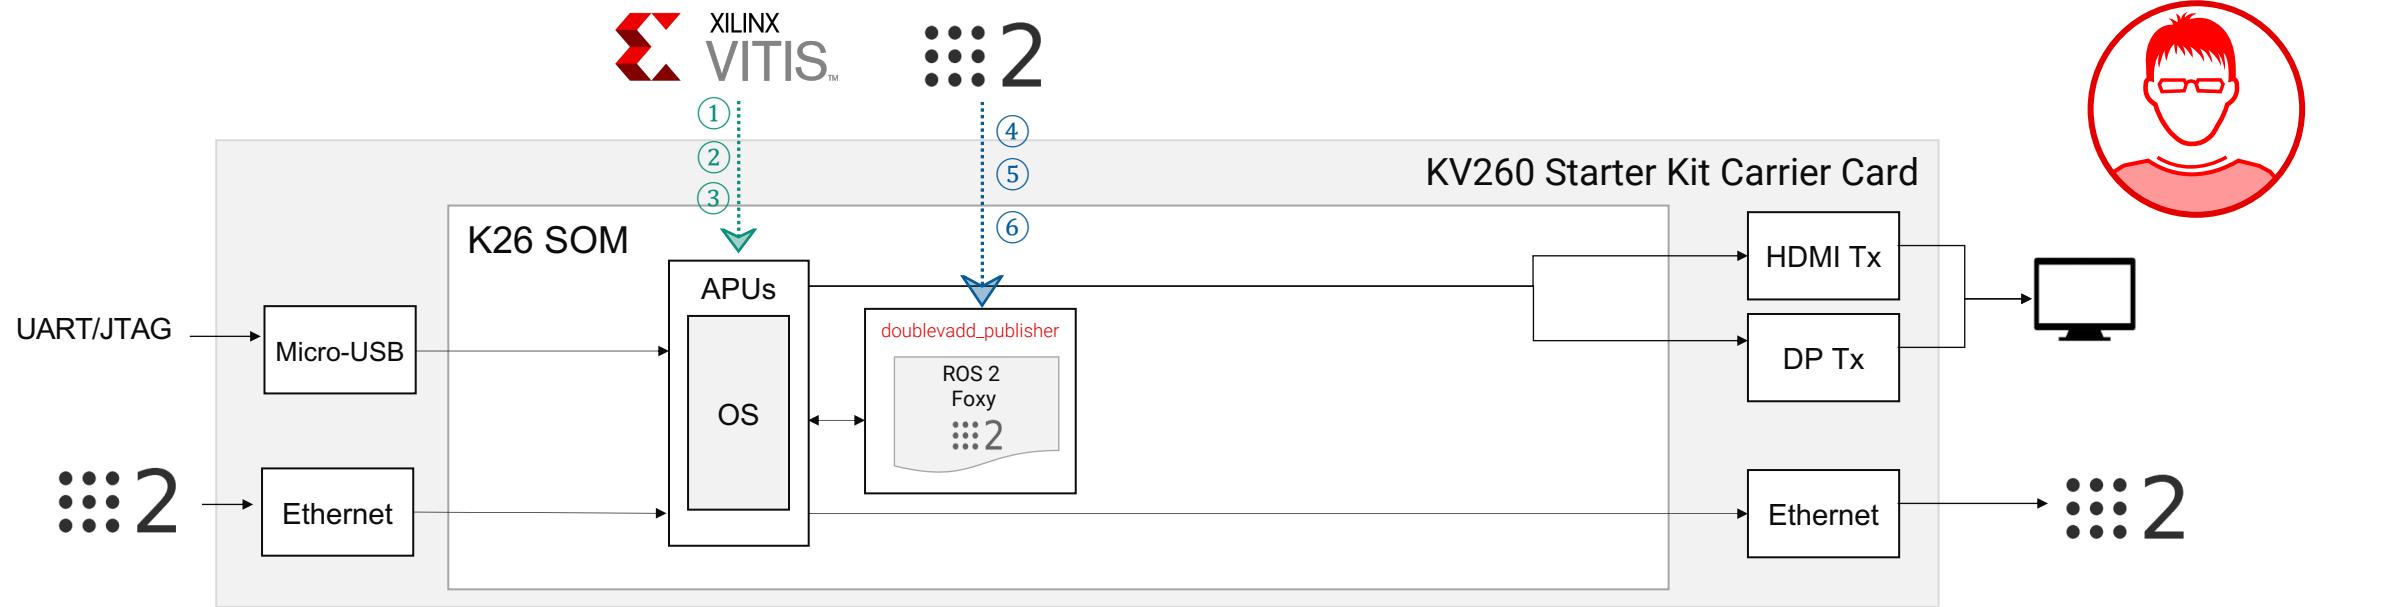
# Accelerating ROS 2

## ROS 2 Hardware acceleration “hello world” – doublevadd\_publisher<sup>1</sup>

<sup>1</sup> [https://github.com/ros-acceleration/acceleration\\_examples/tree/main/doublevadd\\_publisher](https://github.com/ros-acceleration/acceleration_examples/tree/main/doublevadd_publisher)

- ▶ ① Fetch BSP, ② include ROS 2 Yocto recipes, ③ and build rootfs
- ▶ Build ROS 2 doublevadd\_publisher package:
  - Cross-compile ROS 2 workspace: ④ Set up cross-compilation toolchain, ⑤ sysroot, and ⑥ build/cross-compile ROS 2 workspace
  - Yocto recipe: ④ Create Yocto recipe for doublevadd\_publisher, build package
  - Native build: ④ Integrate in rootfs build tools, ⑤ re-build rootfs, and ⑥ build doublevadd\_publisher natively.

Leo “the roboticist”



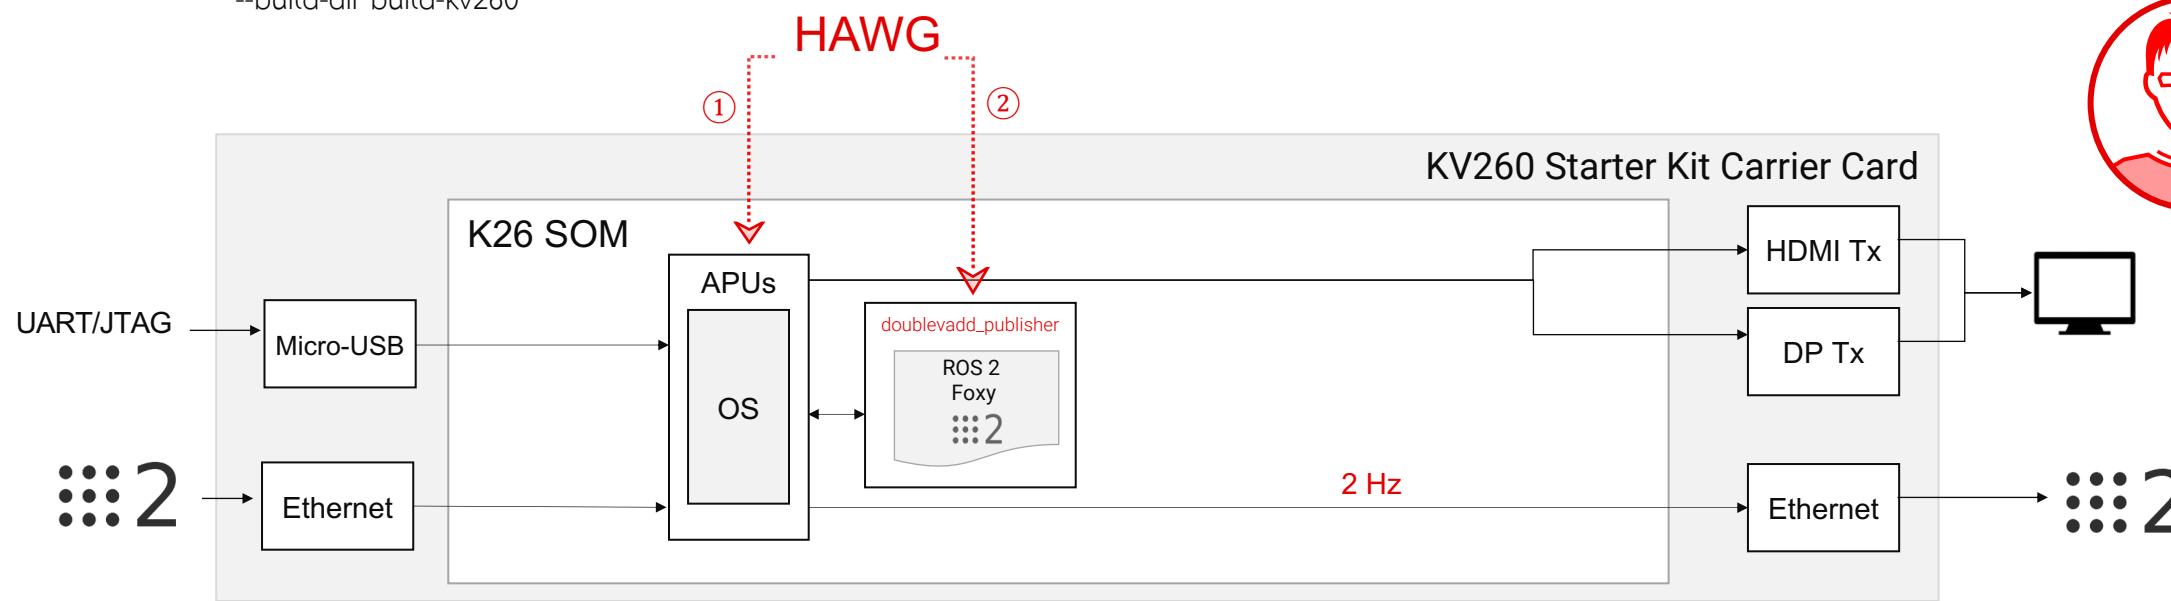
# Accelerating ROS 2 – with HAWG

## ROS 2 Hardware acceleration “hello world” – doublevadd\_publisher<sup>1</sup>

<sup>1</sup> [https://github.com/ros-acceleration/acceleration\\_examples/tree/main/doublevadd\\_publisher](https://github.com/ros-acceleration/acceleration_examples/tree/main/doublevadd_publisher)

# with KRS

```
$ colcon acceleration select kv260          # select KV260 firmware
$ colcon acceleration linux vanilla
$ colcon build --merge-install
$ source install/setup.bash
$ colcon build --merge-install --mixin kv260 \
    --install-dir install-kv260 \
    --build-dir build-kv260                  # ① create roots with ROS 2
                                              # build KRS ROS 2 packages
                                              # apply the overlay
                                              # ② cross-compile doublevadd_publisher
```



Leo “the roboticist”



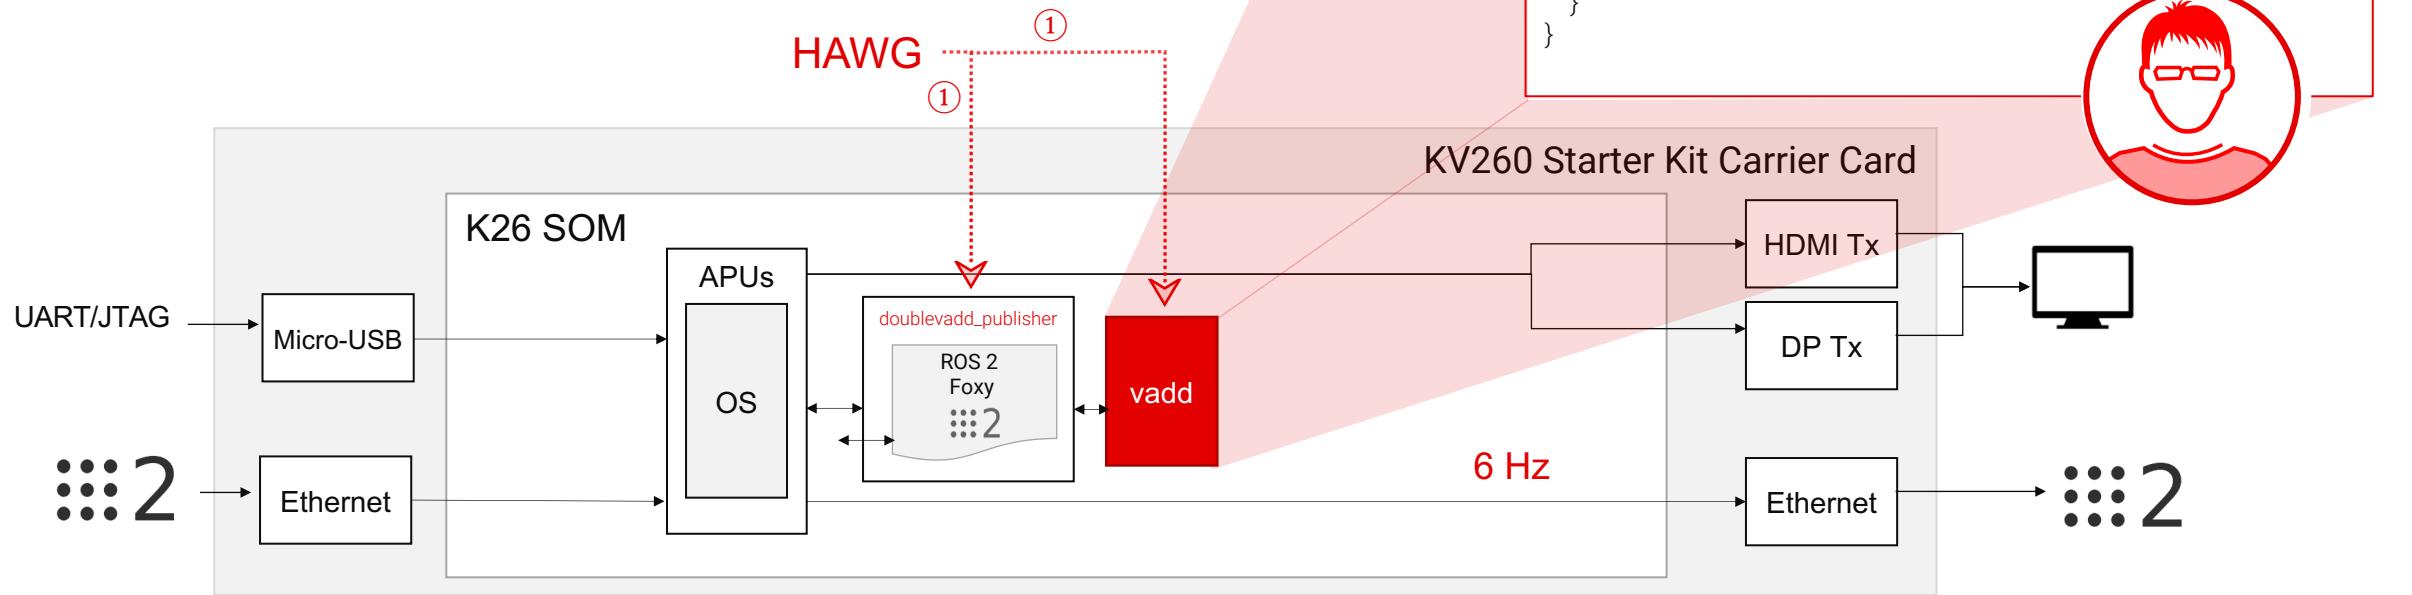
# Accelerating ROS 2 – with HAWG

## ROS 2 Hardware acceleration “hello world”

### – accelerated\_doublevadd\_publisher<sup>1</sup>

<sup>1</sup> [https://github.com/ros-acceleration/acceleration\\_examples/tree/main/accelerated\\_doublevadd\\_publisher](https://github.com/ros-acceleration/acceleration_examples/tree/main/accelerated_doublevadd_publisher)

```
# rebuild ① doublevadd_publisher, including kernels
$ colcon build --merge-install --mixin kv260 \
    --install-dir install-kv260 \
    --build-dir build-kv260
```



# Accelerating ROS 2 – with HAWG

## ROS 2 Hardware acceleration “hello world”

### – faster\_doublevadd\_publisher<sup>1</sup>

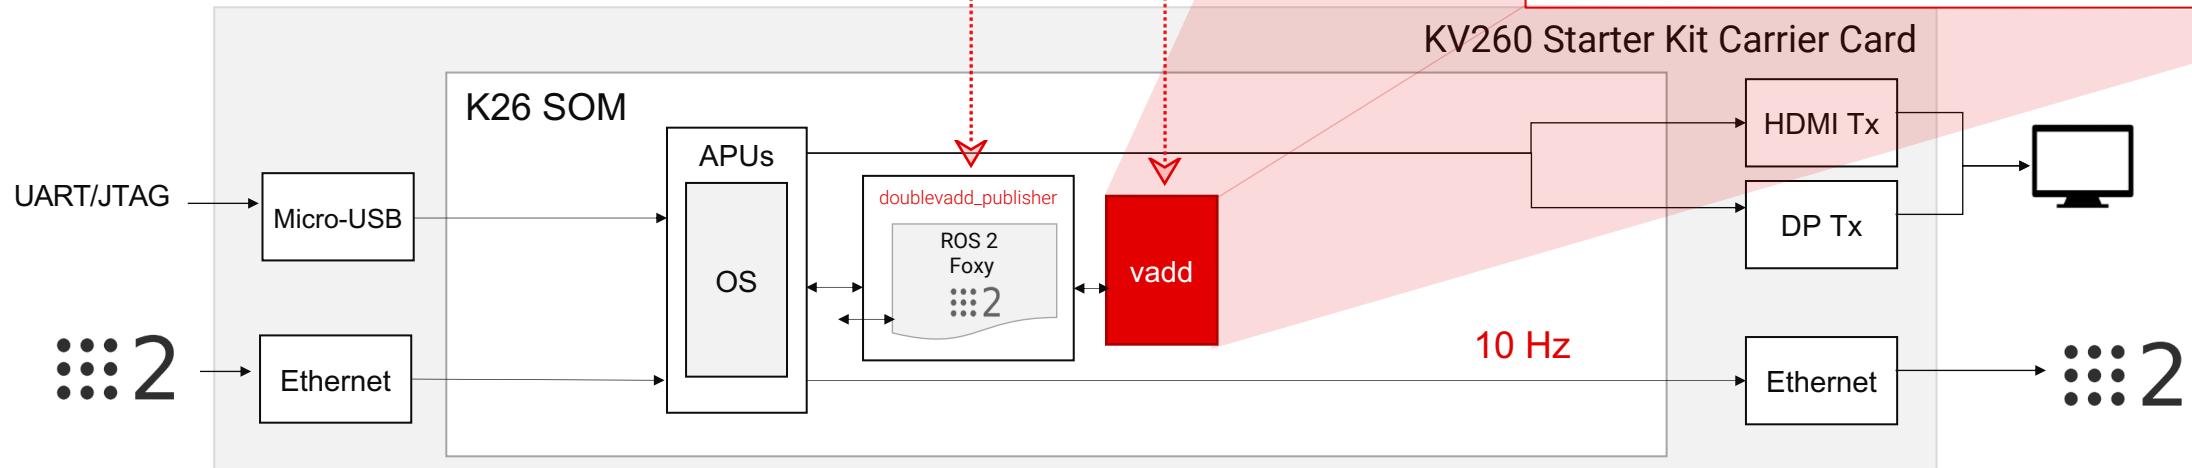
<sup>1</sup> [https://github.com/ros-acceleration/acceleration\\_examples/tree/main/faster\\_doublevadd\\_publisher](https://github.com/ros-acceleration/acceleration_examples/tree/main/faster_doublevadd_publisher)

Leo “the roboticist”



```
# rebuild ① doublevadd_publisher, including kernels
$ colcon build --merge-install --mixin kv260 \
    --install-dir install-kv260 \
    --build-dir build-kv260
```

HAWG  
①



Faster ROS 2 interactions with KRS

```
void vadd(
    const unsigned int *in1, // Read-Only Vector 1
    const unsigned int *in2, // Read-Only Vector 2
    unsigned int *out, // Output Result
    int size // Size in integer
)
{
    #pragma HLS INTERFACE m_axi port=in1 bundle=aximm1
    #pragma HLS INTERFACE m_axi port=in2 bundle=aximm2
    #pragma HLS INTERFACE m_axi port=out bundle=aximm1

    for (int j = 0; j < size; ++j) { // stupidly iterate
        // to generate "double" load
        for (int i = 0; i < size; ++i) {
            for (int i = 0; i < (size/16)*16; ++i) {
                #pragma HLS UNROLL factor=16
                out[i] = in1[i] + in2[i];
            }
        }
    }
}
```

# Accelerating ROS 2 – with KRS

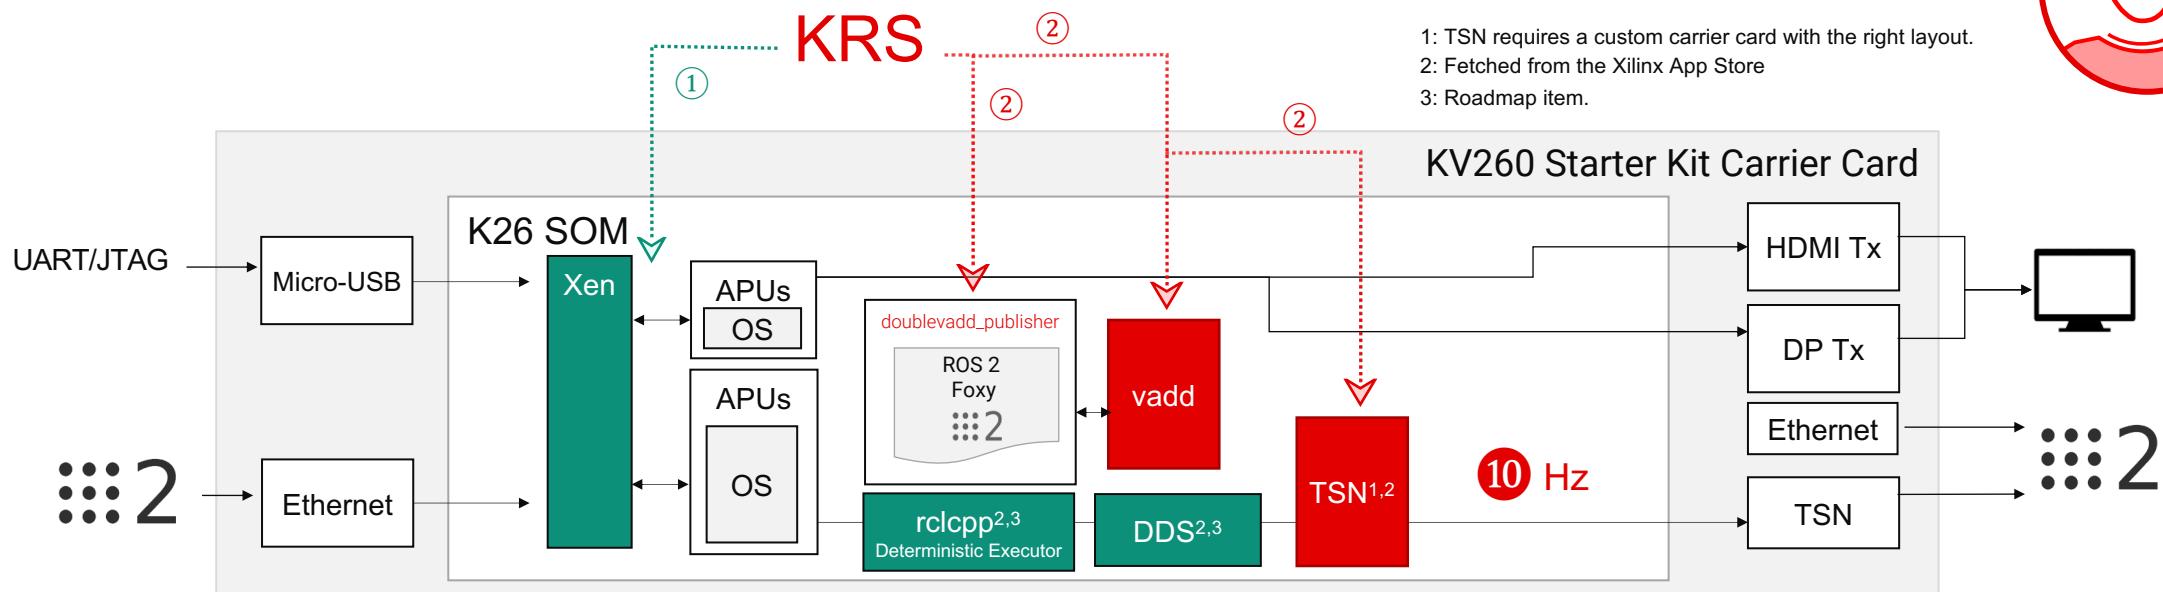
*ROS 2 Hardware acceleration “hello world”*

– doublevadd\_publisher

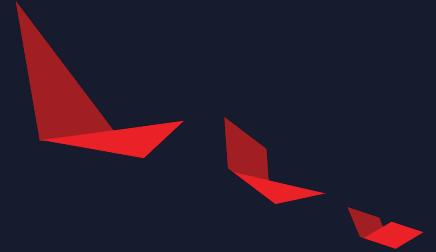
```
# rebuild rootfs for determinism ①
$ colcon acceleration hypervisor \
  --dom0 vanilla \
  --domU preempt_rt
```

```
# rebuild ② doublevadd_publisher, including kernels
$ colcon build --merge-install --mixin kv260 \
  --install-dir install-kv260 \
  --build-dir build-kv260
```

Leo “the roboticist”



Deterministic real-time ROS 2 interactions with KRS with high performance



# ReconROS

# **ReconROS**

## **(ReconOS + ROS)**

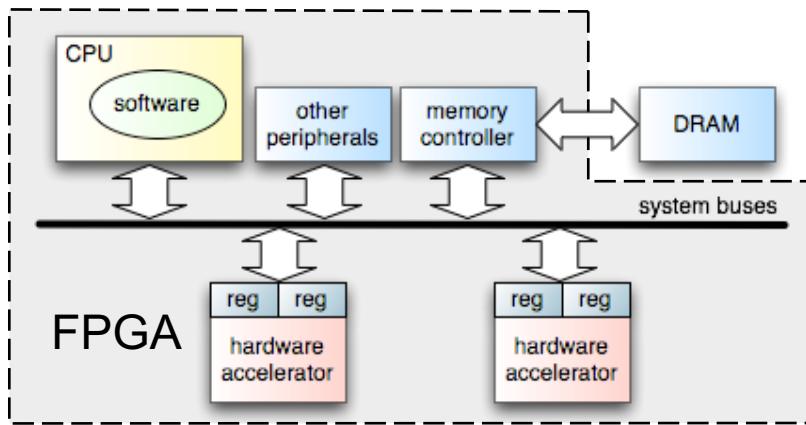
# **Flexible Hardware Acceleration for ROS2 Applications**

Christian Lienen, Marco Platzner  
Paderborn University  
September 29, 2021



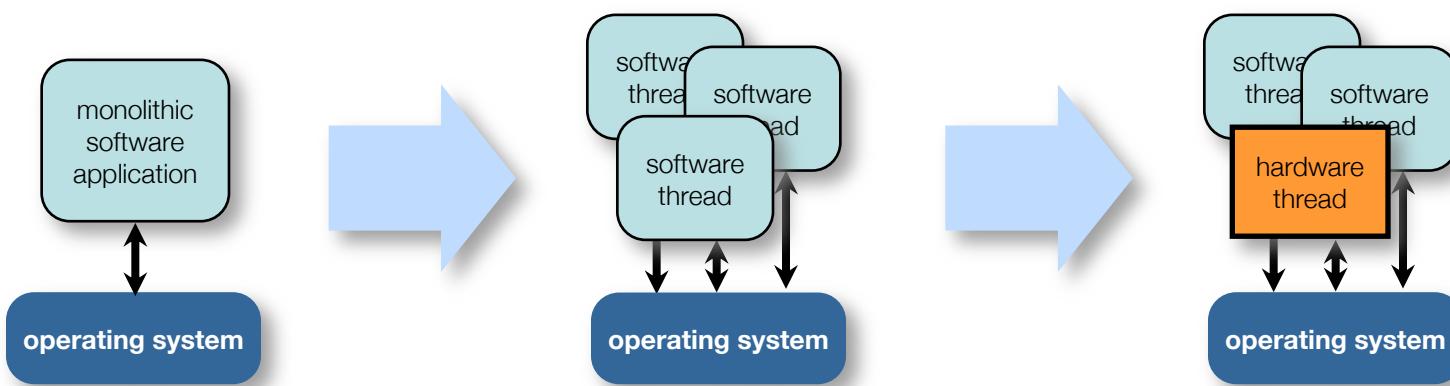
# Operating System Integration of FPGAs

- Traditionally, hardware accelerators integrated as coprocessors
  - software (CPU) explicitly communicates with hardware accelerator
  - memory-mapped registers
  - remote procedure call

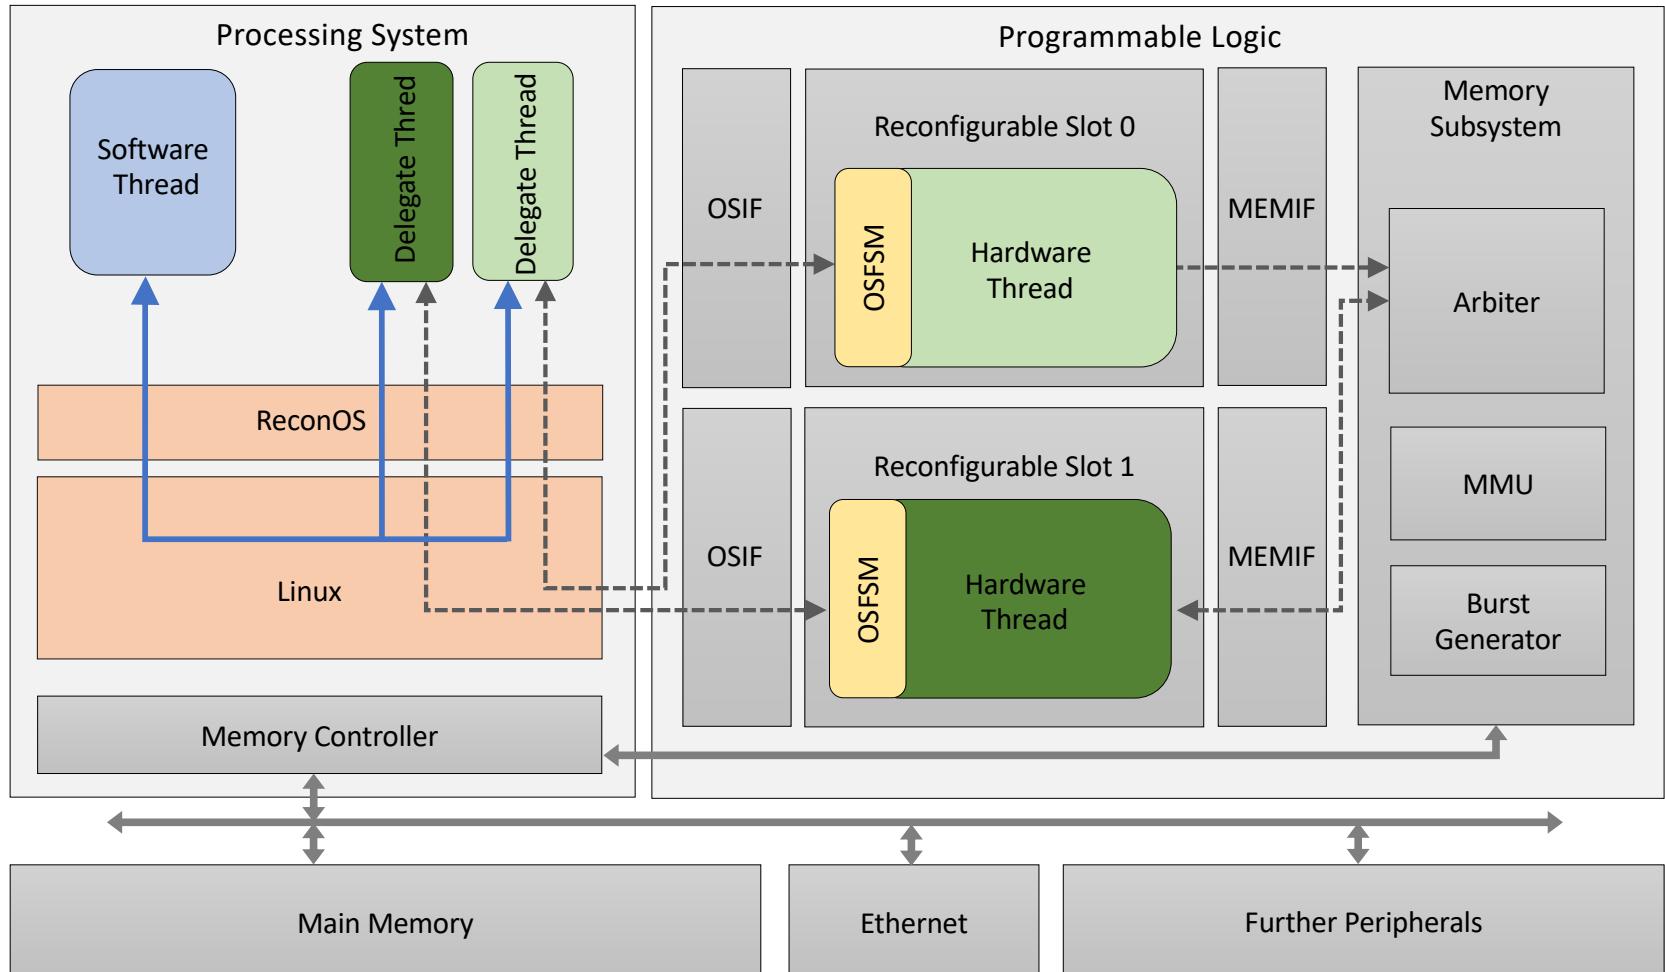


```
int main() {  
    ...  
    data = do_something();  
    io_out32(HW_ACCEL_DATA_REG, data);  
    io_out32(HW_ACCEL_CMD_REG, CMD_GO);  
    while (io_in32(HW_ACCEL_STATUS_REG) != DONE)  
        sleep();  
    result = io_in32(HW_ACCEL_RESULT_REG);  
    ...  
}
```

- Extend multithreaded programming model to hardware
  - hardware accelerators are turned into hardware threads
  - threads communicate and synchronize using programming model primitives, e.g., semaphores, mutexes, mailboxes, shared memory read/write
  - established software abstraction (e.g., POSIX pthreads) provides a standarized interface between software and hardware



# Example ReconOS Hardware Architecture



- Hardware threads use cooperative multitasking
- Hardware threads can be loaded / removed by partial reconfiguration

# ReconOS Versions

## Version 1.0

eCos / PowerPC

Virtex-2, Virtex-2Pro, Virtex-4

## Version 2.0

Linux, eCos / PowerPC

Virtex-2Pro, Virtex-4

virtual memory support

## Version 3.0

Linux, xilkernel / MicroBlaze

Virtex-6

## Version 3.1

Linux / ARM

Xilinx Zynq

ISE

## Version 4.0

Linux / ARM

Xilinx Zynq

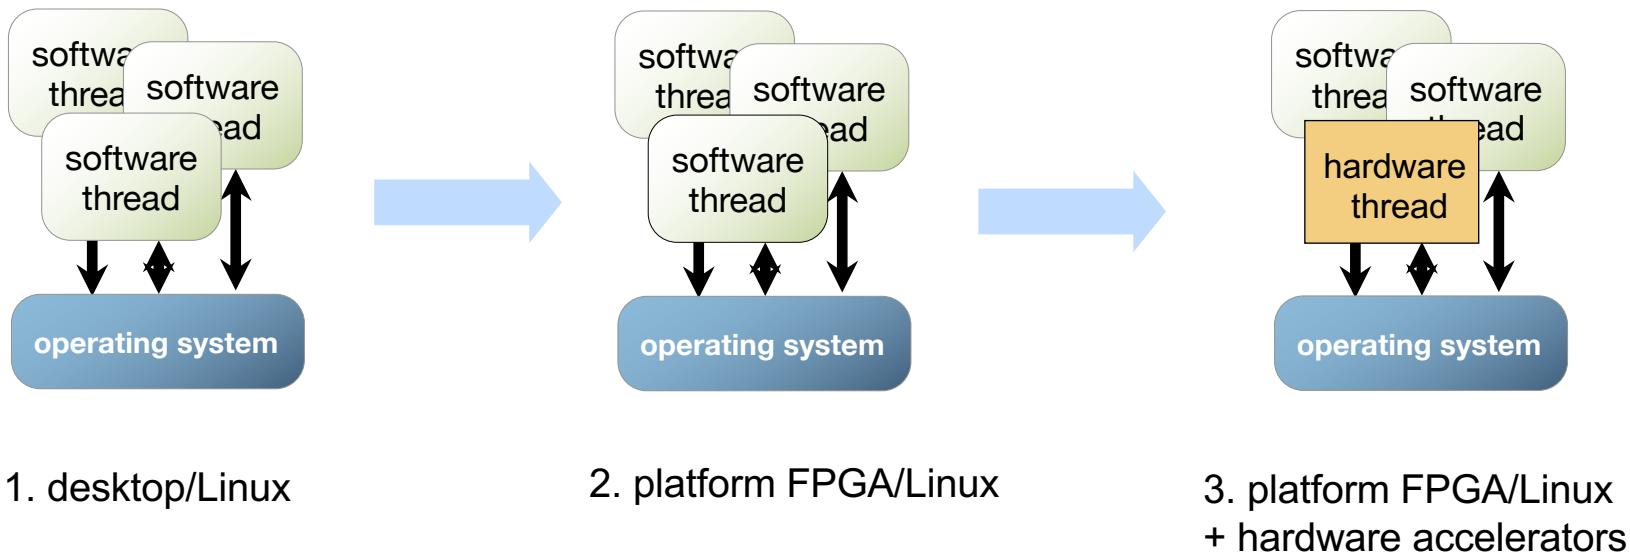
Vivado, HLS for hardware thread design,  
flexible and dynamic clocks

The screenshot shows the official website for ReconOS. At the top, there's a navigation bar with links to "About ReconOS", "Getting Started", "Documentation", "Get Involved", and "Publications". Below the navigation is a large banner image showing a close-up of a green printed circuit board (PCB) with a central black integrated circuit chip. Overlaid on the banner is the text "Support for the Xilinx Zynq" and "ReconOS now officially supports the Xilinx Zynq Platform". Below the banner, there's a quote: "The ReconOS operating system for reconfigurable computing offers a unified multithreaded programming model and OS services for threads executing in software and threads mapped to reconfigurable hardware. By semantically integrating hardware accelerators into a standard OS environment, ReconOS allows for rapid design-space exploration, supports a structured application development process, and improves the portability of applications between different reconfigurable computing systems." Below the quote is the tagline "ReconOS – an operating system approach for reconfigurable computing". The main content area features three sections with icons: "Multithreaded Programming Model" (represented by a hand icon), "Active development and support" (represented by a person icon), and "Extended and easy to use Toolchain" (represented by a wrench icon). Each section includes a brief description. At the bottom of the page, there's a footer with logos for Universität Paderborn and EPICS, and a prominent blue button with the website address "www.reconos.de".

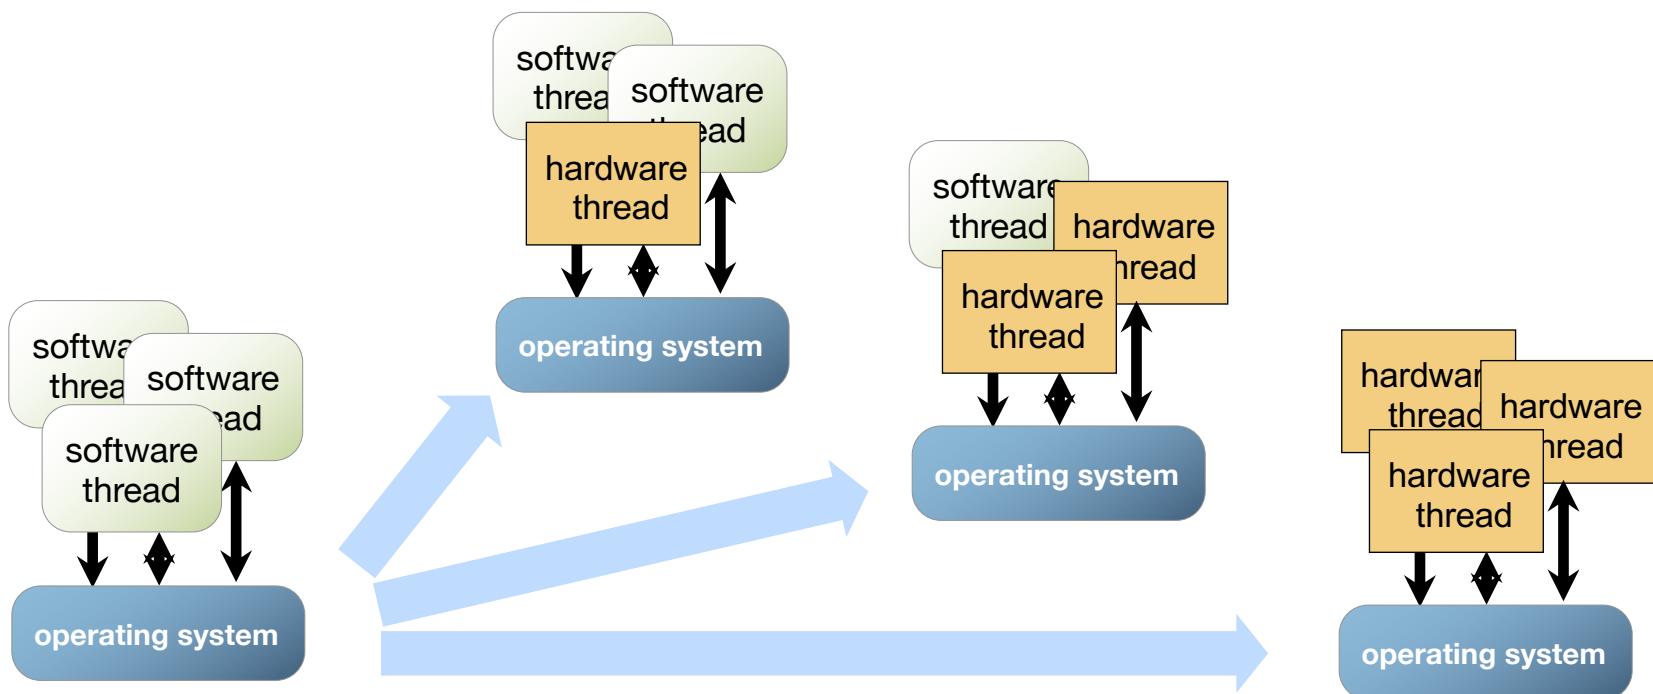


# Experience with ReconOS

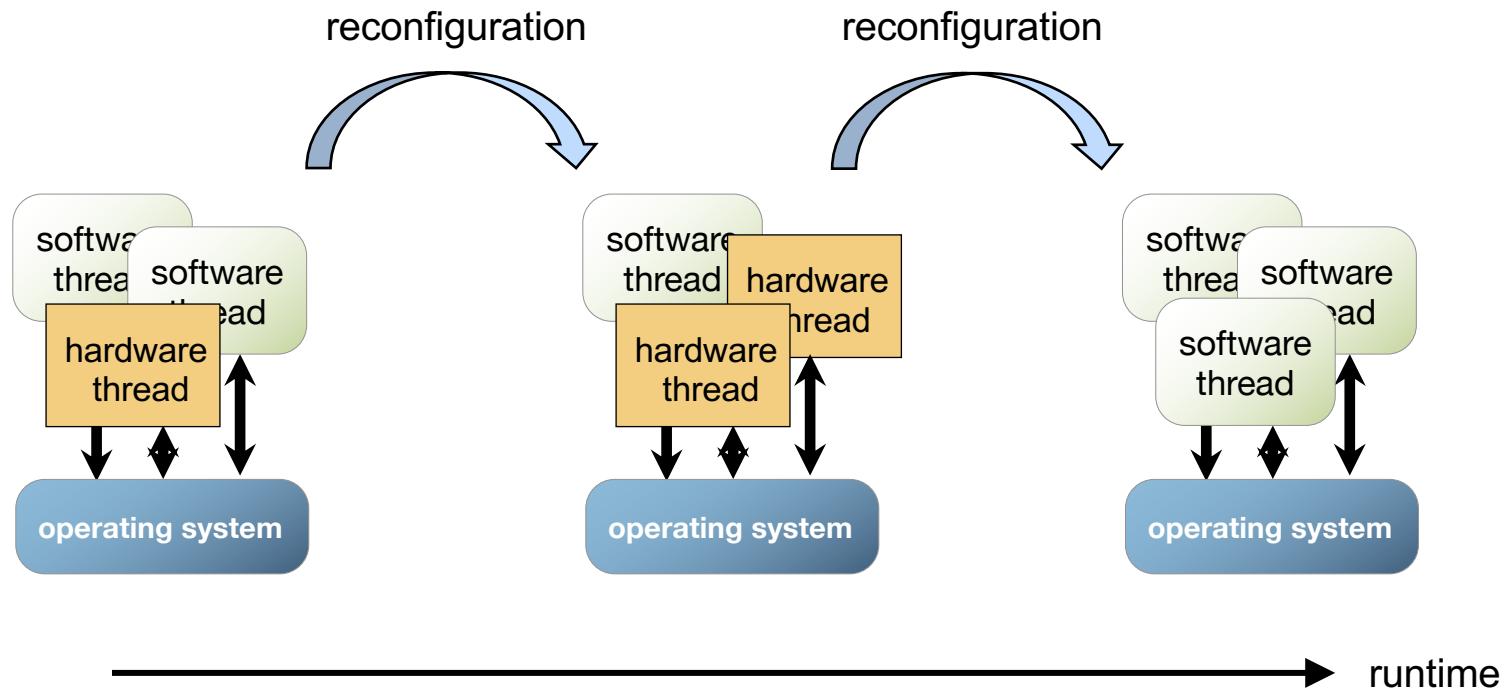
1. ReconOS supports a step-by-step application design process



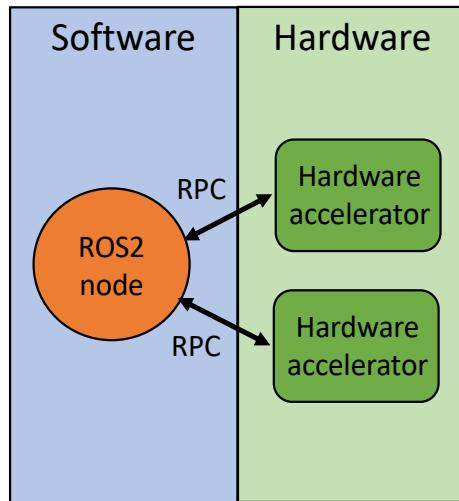
## 2. ReconOS facilitates fast design space exploration



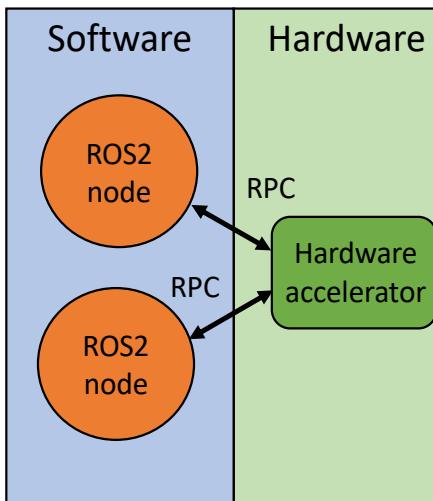
- ReconOS enables migration between software and hardware at runtime → (self-)adaptive systems



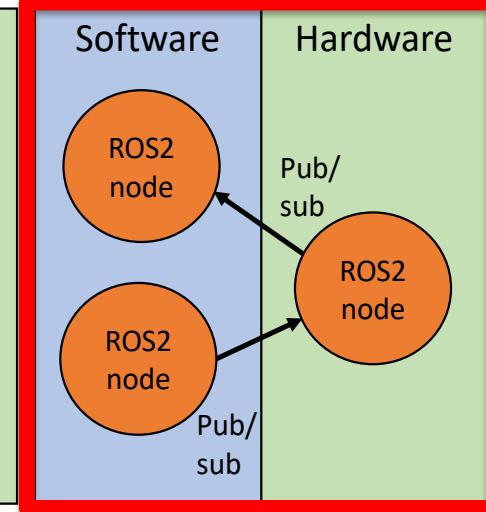
# Hardware Acceleration Options for ROS Nodes



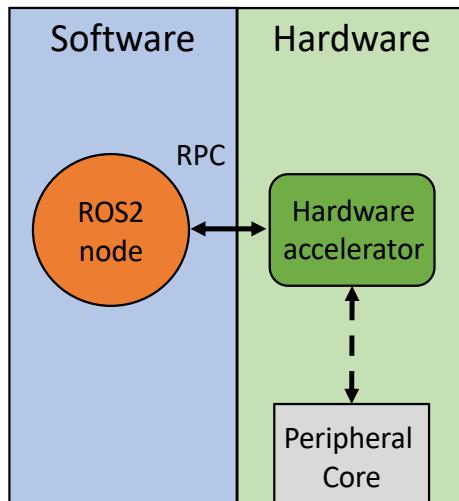
(a)



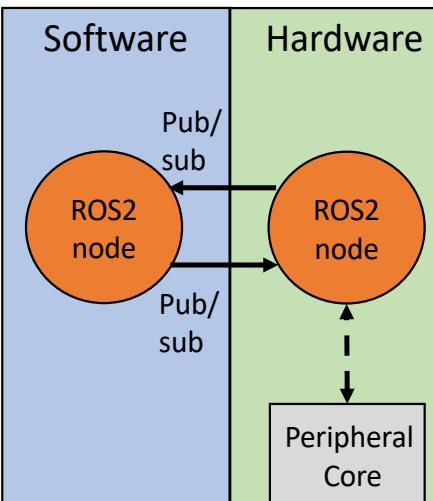
(b)



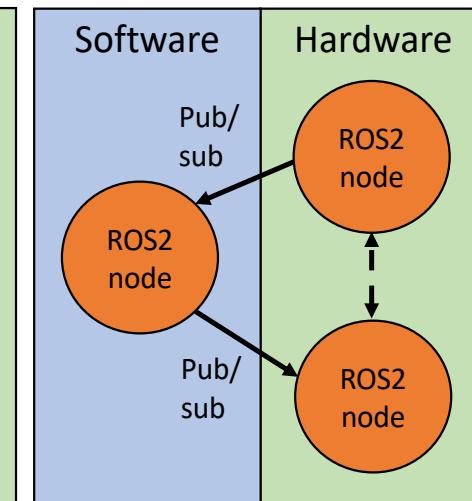
(c)



(d)

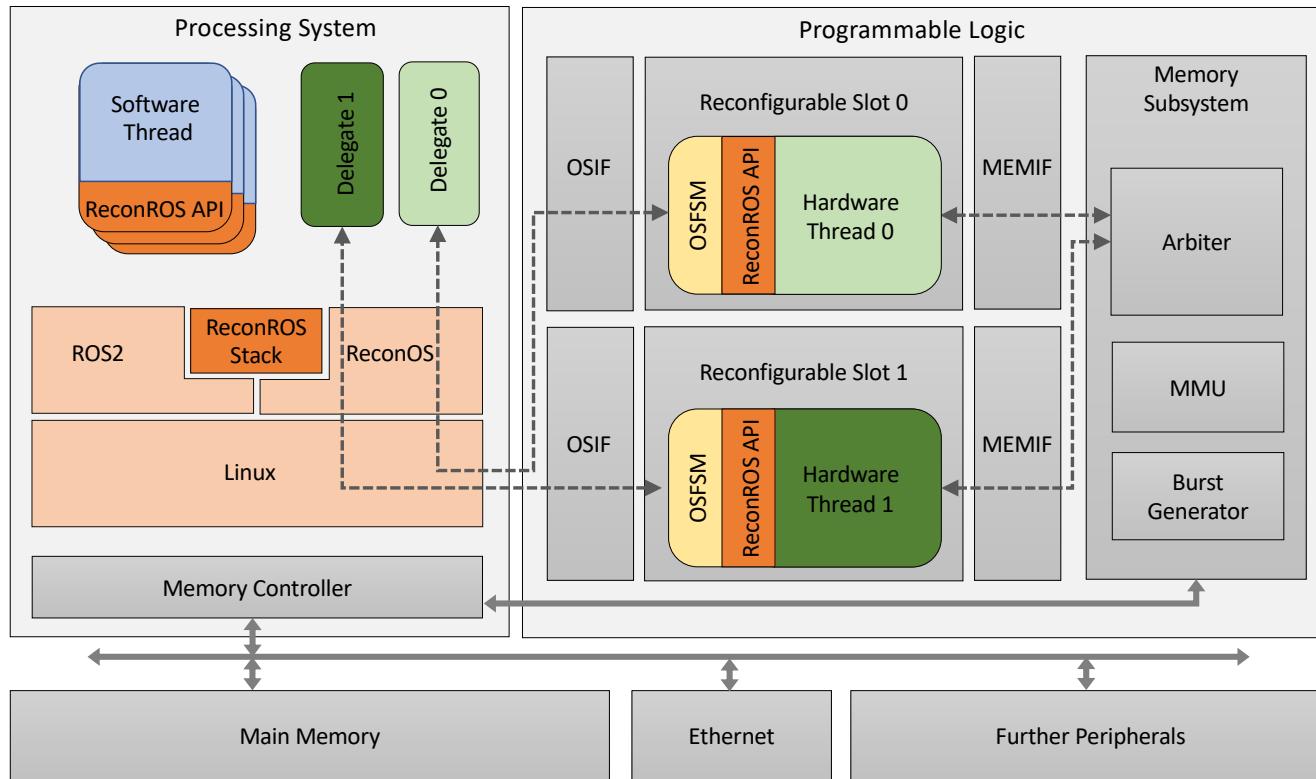


(e)



(f)

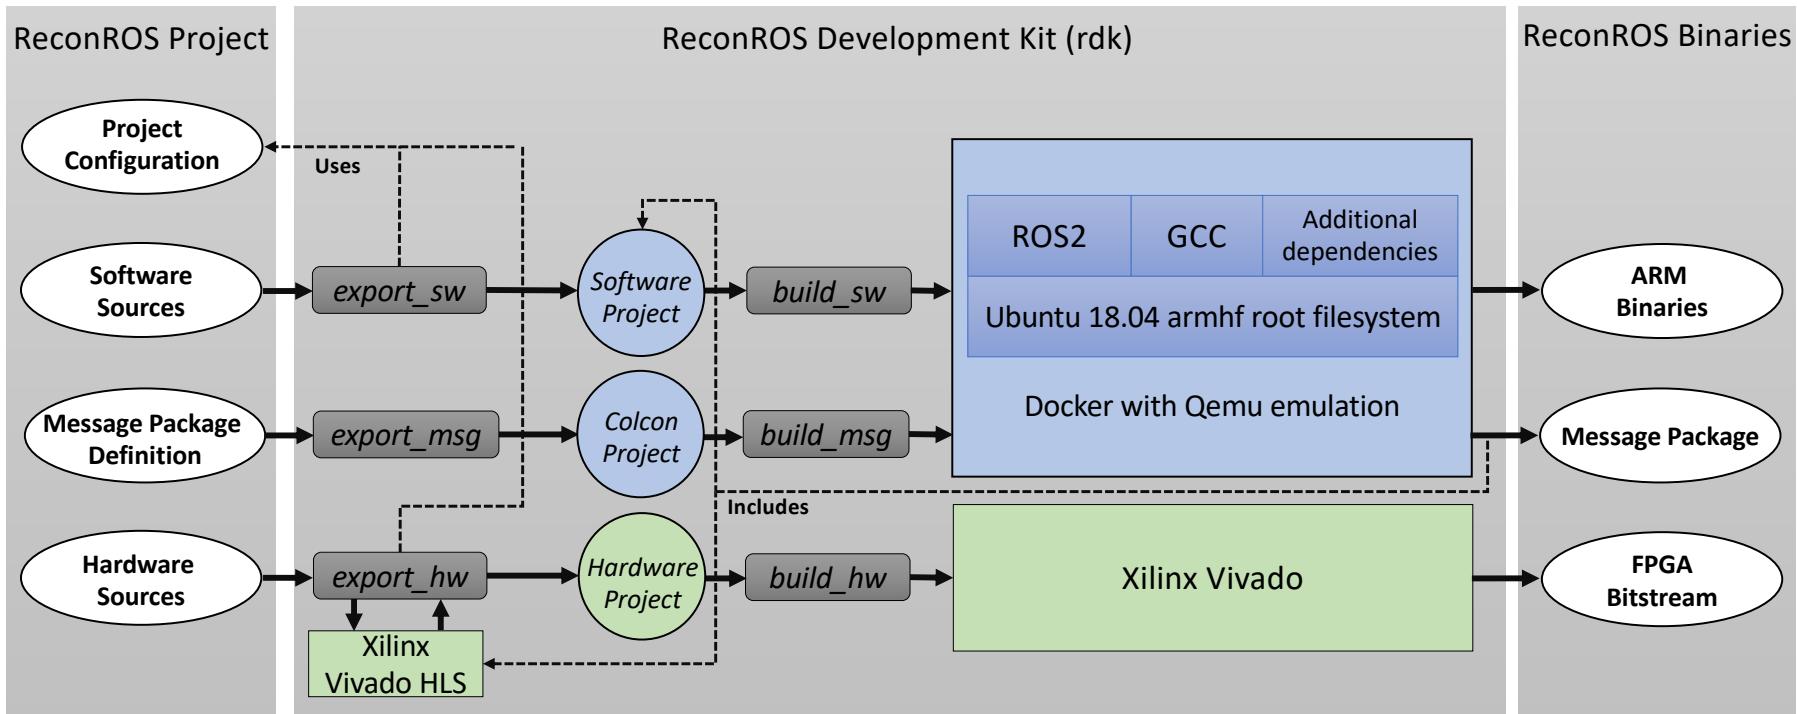
# ReconROS Architecture



- **ReconROS API**
  - provides access to ReconROS stack for hardware and software threads
  - easily extendable and slightly abstracted subset of ROS 2 – *rcl* API
- **ReconROS stack**
  - ROS 2-related objects added as ReconOS primitives (e.g. ROS 2 nodes, pub/sub, services, actions...)
  - stack includes initializations and memory management



# ReconROS Design Flow

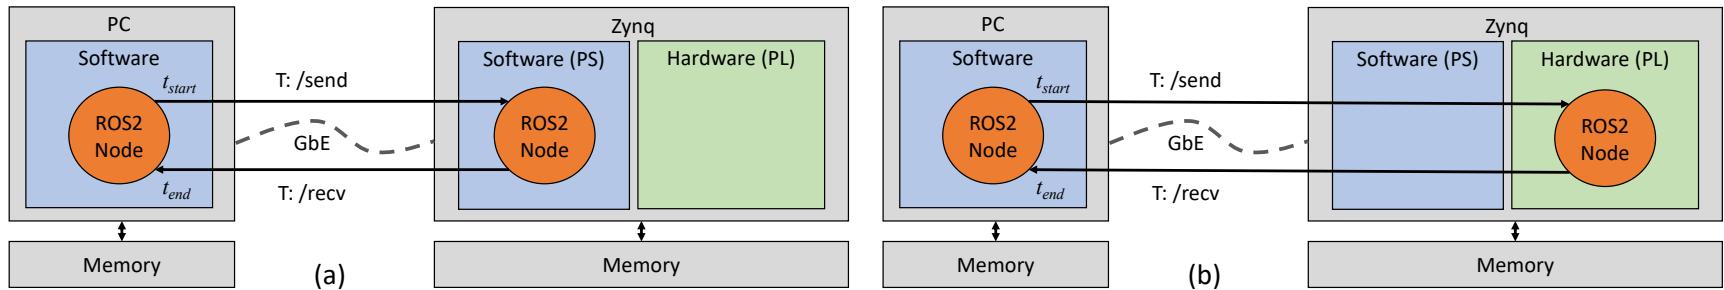


- **Software / message tool chain**
  - support of C/C++ source files and ROS 2 message definition files
  - uses cross compilation to ARM for easier handling of dependencies
- **Hardware tool chain**
  - support of C/C++ HLS sources and VHDL / Verilog descriptions
  - message definitions available in VIVADO HLS as C structs

# Benefits of Mapping ROS Nodes to Hardware

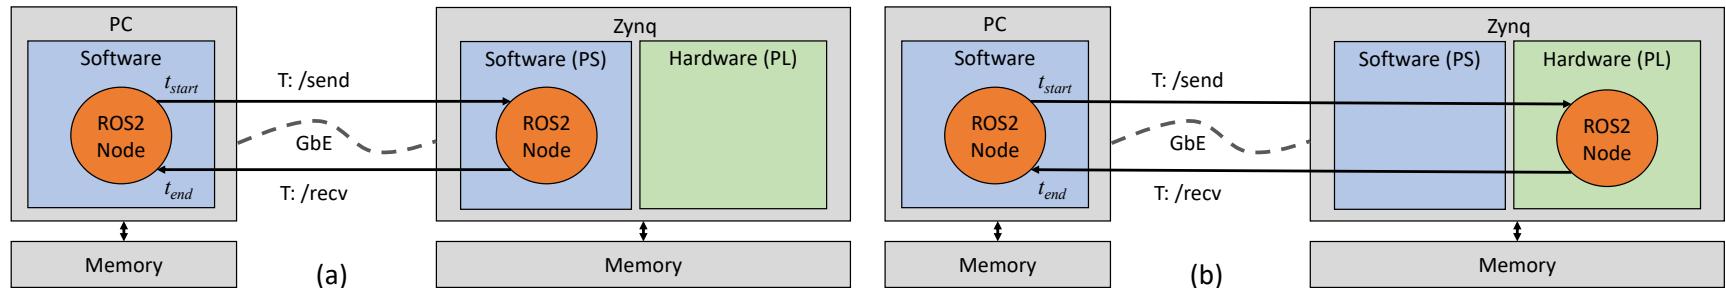
- Higher performance
  - raw speedup depends on node function and the effort spent for optimization
  - overall speedup depends also on overall ROS application (Amdahl's law)
  - hardware nodes execute in parallel to CPU cores → lower CPU load and/or higher performance for overall application
- Potentially lower jitter
- Power/energy savings
- Pre/post-processing in hardware
  - when sensor/actuator is directly connected to FPGA





Message size	$t_{pp\text{-}echo-SW}$ [ms]	$t_{pp\text{-}echo-HW}$ [ms]	$S_{pp\text{-}echo}$
4 Byte	0.81	1.2	0.68
8 KiB	10.65	10.48	1.02
1 MiB	52.21	52.16	1.01
6 MiB	363.91	363.30	1.00
10 MiB	630.37	624.02	1.01

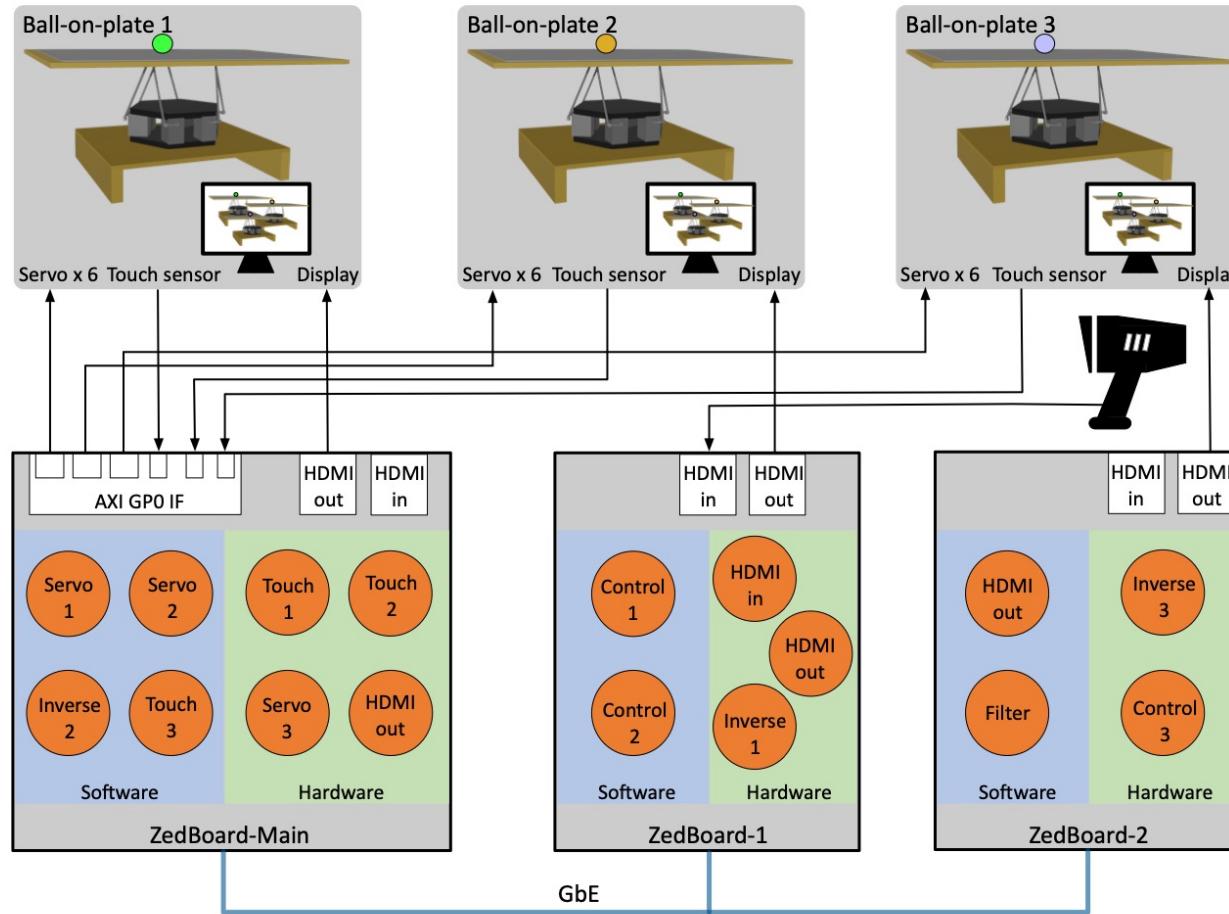
- ‘Echo’ application
  - ROS nodes just pass message pointers between subscribe() and publish(), but there is no real computation
  - no overhead for using hardware (underlying ReconOS signaling)



RECONROS application	Message Size In/Out	$t_{raw-SW}$ [ms]	$t_{raw-HW}$ [ms]	$S_{raw}$	$t_{pp-SW}$ [ms]	$t_{pp-HW}$ [ms]	$S_{pp}$
Inverse	4/4 Byte	1.20	0.19	6.32	7.70	6.64	1.16
Sorting	8/8 KiB	17.44	35.11	0.50	24.42	42.08	0.58
Sobel	900/900 KiB	37.53	22.28	1.68	83.39	68.54	1.22
MNIST	850/4 Byte	88.03	30.74	2.86	98.58	41.25	2.39

- ‘Smaller’ applications
  - hardware nodes created with high-level synthesis (HLS) from C/C++ without any optimization

- Mechatronics application
  - easy to create new system configurations with different distributions of nodes between compute boards and between software and hardware



# Summary, Sources and Literature

ReconROS enables ROS 2 application developers to use flexible hardware acceleration with FPGAs through a consistent programming model across the hardware/software boundary.

ReconOS



<https://github.com/ReconOS/reconos>

- [1] Enno Lübers and Marco Platzner. "[ReconOS: "Multithreaded programming for reconfigurable computers"](#)", ACM Transactions on Embedded Computing Systems, 9(1):1-33, 2009.
- [2] Andreas Agne, Markus Happe, Ariane Keller, Enno Lübers, Bernhard Plattner, Marco Platzner and Christian Plessl. "[ReconOS: An Operating System Approach for Reconfigurable Computing](#)", IEEE Micro, 34(1):60-71, 2014.

ReconROS



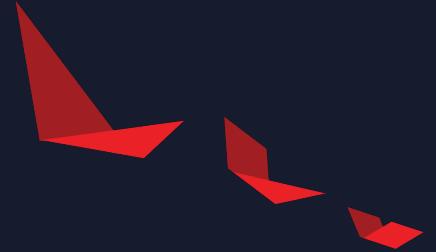
<https://github.com/Lien182/ReconROS>

(Code, Demos, Board Images,...)

- [3] Christian Lienen, Marco Platzner and Bernhard Rinner. "[ReconROS: Flexible Hardware Acceleration for ROS2 Applications](#)", International Conference on Field Programmable Technology (ICFPT), 2020.

- [4] Christian Lienen and Marco Platzner. "[Design of Distributed Reconfigurable Robotics Systems with ReconOS](#)", ArXiv paper under review, 2021.





# Q&A